Selected measured outcomes

Four experiments. Four separate comparisons.

01 / BIN PACKING

Training excess · lower is better

4.590% → 1.829%

1.379% on private split · 60 model calls

One recorded trajectory

02 / CPU KERNEL

Private speedup

16.06×

20 model calls · same-binary timing

Naive baseline, not optimized BLAS

03 / VEHICLE ROUTING

Private gap · score-only versus memory

0.748% → 0.581%

One seed · shared ancestry and memory

Timing noise; memory effect not isolated

04 / VMEC++ POLISH

Official P2 gain

+0.001956

17 steps · 715 solver calls

Inherited design · local gain, then stall

Different tasks, budgets and units · Evidence review: 2026-09-17