


From physics measurements to selection and feedback

A fresh evaluator re-scores the returned boundary; the official benchmark score stays separate from search-only shaping and diagnostics.

A Clean-room score contract

Candidate code is absent from verification



B Selection pressure

Applied to non-private splits after fresh evaluation

$$S_{\text{select}} = S_{\text{base}} - P_{\text{margin}} - P_{\text{novelty}}$$

FEASIBILITY-MARGIN PENALTY

$$P_{\text{margin}} = 0.92 \max(0, f - 0.002)$$

Only when P2 > 0; reported as margin_penalty + honest_score.

PUBLIC-BANK NOVELTY PENALTY

$$P_{\text{novelty}} = 0.05 (1 - d/0.001) \text{ if } d < 0.001; \text{ otherwise } 0$$

Applied only with a positive score and available same-NFP distance d ; penalty is rounded to 4 decimals.

Infeasible: penalties inactive $\Rightarrow$ selection = $-f$.

C Returned feedback

EvalResult → ledger → selected feedback strategy

HARNESS VALUE

score · error · seconds

CORE PHYSICS

p2_score · feasibility · objective_L

NON-PRIVATE DIAGNOSTICS

aspect_ratio · max_elongation · edge_iota_per_nfp
edge_mirror_ratio · log10_qi · shaped_score

CONDITIONAL / ACCOUNTING

bank_dist · bank_cos · penalty fields
evals_used · seconds_used · submittable

MODE

score only

reflections

memory

Official P2 the fresh pinned evaluator's benchmark score is returned unchanged; honest_score and tolerance_used are diagnostics only.

Scope Penalty constants are environment-overrideable and fingerprinted in experiment_spec; the diagram shows code defaults, not measured effects.